

Recommended Exercises (Module 2)

Jim Totland

1/16/2021

[Link to problem set](#)

Problem 1

- Weather forecasting. Response: “Sunny”, “Cloudy”, “Rain” etc. Predictors: Air pressure, temperature, and the weather of the previous day(s). The goal is to predict.
- Battery life of a phone. Response: Time until the phone is dead. Predictors: Screen size, Battery specs, Processor etc. Both prediction and inference are relevant here. Given a phone, we want to be able to predict what the battery life will be, based to the predictors, but from the regression we will also be able to infer which predictors are most significant.

Problem 2

- In this example, the more flexible methods have a smaller test MSE. But at some point the test MSE start to increase monotonically with the flexibility. This is a result of overfitting.
- The variance refers to how much \hat{f} would change if we used another set of training data. A small variance could indicate that a rigid method has been used, implying that the data is most likely underfitted.
- Bias generally decreases with flexibility, which indicates that a very low bias is connected to overfitting the data.

Problem 3

```
library(ISLR)
data(Auto)
```

- Use the `glimpse` function from the tidyverse:

```
glimpse(Auto)
```

```
## Rows: 392
## Columns: 9
## $ mpg      <dbl> 18, 15, 18, 16, 17, 15, 14, 14, 14, 15, 15, 14, 15, 14...
## $ cylinders <dbl> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 4, 6, 6, 6, ...
## $ displacement <dbl> 307, 350, 318, 304, 302, 429, 454, 440, 455, 390, 383,...
## $ horsepower <dbl> 130, 165, 150, 150, 140, 198, 220, 215, 225, 190, 170,...
## $ weight     <dbl> 3504, 3693, 3436, 3433, 3449, 4341, 4354, 4312, 4425, ...
## $ acceleration <dbl> 12.0, 11.5, 11.0, 12.0, 10.5, 10.0, 9.0, 8.5, 10.0, 8....
## $ year       <dbl> 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70, 70...
## $ origin     <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 3, 1, 1, 1, ...
## $ name       <fct> chevrolet chevelle malibu, buick skylark 320, plymouth...
```

The data has dimensions 392×9 . All predictors except `name` are quantitative, although some of them may also be treated as categorical.

b) The range is found by applying the `range()` function. For example:

```
range(Auto$mpg)
```

```
## [1] 9.0 46.6
```

Alternatively use `sapply`:

```
quant = c(1,3,4,5,6,7)
sapply(Auto[, quant], range)
```

```
##      mpg displacement horsepower weight acceleration year
## [1,] 9.0           68          46   1613           8.0   70
## [2,] 46.6         455         230   5140          24.8   82
```

c) The mean and standard deviation can be found in the following way:

```
for (i in 1:8) {
  print(summarise(Auto, mean = mean(Auto[,i]), sd = sd(Auto[,i])))
}
```

```
##      mean      sd
## 1 23.44592 7.805007
##      mean      sd
## 1 5.471939 1.705783
##      mean      sd
## 1 194.412 104.644
##      mean      sd
## 1 104.4694 38.49116
##      mean      sd
## 1 2977.584 849.4026
##      mean      sd
## 1 15.54133 2.758864
##      mean      sd
## 1 75.97959 3.683737
##      mean      sd
## 1 1.576531 0.8055182
```

d) Possible, though not very clean, solution:

```
ReducedAuto <- Auto[- (10:85),]
```

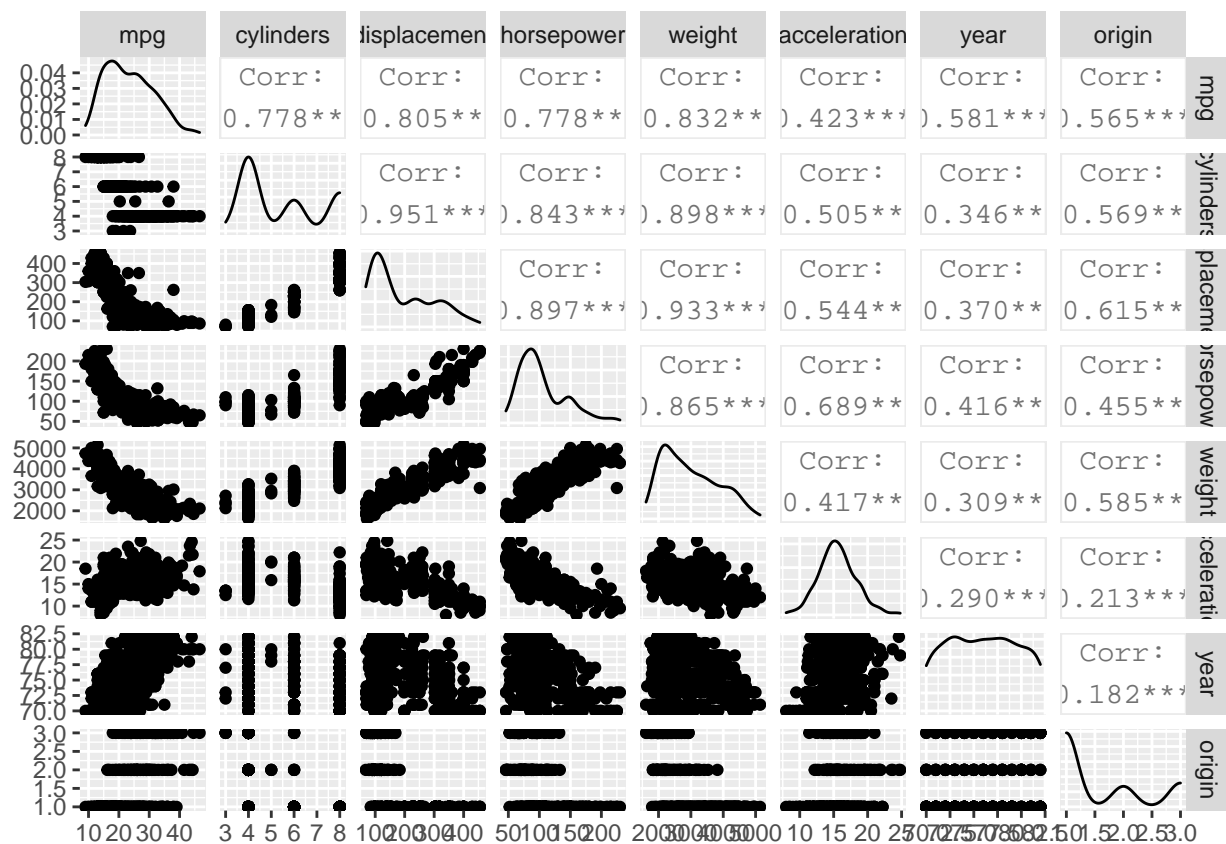
```
for (i in 1:8) {
  print(summarise(ReducedAuto, mean = mean(ReducedAuto[,i]),
    sd = sd(ReducedAuto[,i]),
    range = range(ReducedAuto[,i])))
}
```

```
##      mean      sd range
## 1 24.40443 7.867283 11.0
## 2 24.40443 7.867283 46.6
##      mean      sd range
## 1 5.373418 1.654179 3
## 2 5.373418 1.654179 8
##      mean      sd range
## 1 187.2405 99.67837 68
```

```
## 2 187.2405 99.67837 455
##      mean      sd range
## 1 100.7215 35.70885 46
## 2 100.7215 35.70885 230
##      mean      sd range
## 1 2935.972 811.3002 1649
## 2 2935.972 811.3002 4997
##      mean      sd range
## 1 15.7269 2.693721 8.5
## 2 15.7269 2.693721 24.8
##      mean      sd range
## 1 77.14557 3.106217 70
## 2 77.14557 3.106217 82
##      mean      sd range
## 1 1.601266 0.81991 1
## 2 1.601266 0.81991 3
```



e)

```
library(GGally)
ggpairs(Auto[-9])
```



From the plot we can see that there seems to be a linear relationship between multiple predictors. E.g. **weight** and **displacement** have a clearly positive linear relationship. There also seems to be some non-linear relationships, e.g. between **mpg** and **horsepower**.

f) I will here treat **cylinders** and **origin** as qualitative variables and get the following box plots:



The majority of the variables seem to have some relevance in predicting mpg. But the variables year,

acceleration and name are probably the least impactful based on visual inspection.

g) The following function calculates the correlation matrix given the covariance matrix.

```
getCor <- function(covMat) {  
  rows <- dim(covMat)[1]  
  cols <- dim(covMat)[2]  
  corMat <- matrix(nrow = rows, ncol = cols)  
  
  for (i in 1:rows) {  
    for(j in 1:cols) {  
      corMat[i,j] = covMat[i,j] / (sqrt(covMat[i,i]) * sqrt(covMat[j,j]))  
    }  
  }  
  return (corMat)  
}
```

Problem 4


a)

```
library(MASS)  
  
##  
## Attaching package: 'MASS'  
## The following object is masked from 'package:dplyr':  
##  
##      select  
v1 <- as.data.frame(mvrnorm(n = 1000,  
  mu = c(2,3),  
  Sigma = matrix(c(1, 0, 0, 1), 2, 2, byrow = T) ))  
colnames(v1) = c("x1", "x2")  
  
v2 <- as.data.frame(mvrnorm(n = 1000,  
  mu = c(2,3),  
  Sigma = matrix(c(1, 0, 0, 5), 2, 2, byrow = T) ))  
colnames(v2) = c("x1", "x2")  
  
v3 <- as.data.frame(mvrnorm(n = 1000,  
  mu = c(2,3),  
  Sigma = matrix(c(1, 2, 2, 5), 2, 2, byrow = T) ))  
colnames(v3) = c("x1", "x2")  
  
v4 <- as.data.frame(mvrnorm(n = 1000,  
  mu = c(2,3),  
  Sigma = matrix(c(1, -2, -2, 5), 2, 2, byrow = T) ) )  
colnames(v4) = c("x1", "x2")
```

b) Plot the simulated distributions:

```
library(gridExtra)  
  
##  
## Attaching package: 'gridExtra'  
## The following object is masked from 'package:dplyr':
```

```
##
##      combine
p1 <- ggplot(v1, aes(x1, x2)) + geom_point() + labs(title = "set1") + theme_minimal()
p2 <- ggplot(v2, aes(x1, x2)) + geom_point() + labs(title = "set2") + theme_minimal()
p3 <- ggplot(v3, aes(x1, x2)) + geom_point() + labs(title = "set3") + theme_minimal()
p4 <- ggplot(v4, aes(x1, x2)) + geom_point() + labs(title = "set4") + theme_minimal()
grid.arrange(p1,p2,p3,p4, ncol = 2)
```



Problem 5

a) Supplied code:

```
library(ggplot2)
library(ggpubr)
set.seed(2) # to reproduce
M = 100 # repeated samplings, x fixed
nord = 20 # order of polynoms
x = seq(from = -2, to = 4, by = 0.1)
truefunc = function(x) {
  return(x^2)
}
true_y = truefunc(x)
error = matrix(rnorm(length(x) * M, mean = 0, sd = 2), nrow = M, byrow = TRUE)
ymat = matrix(rep(true_y, M), byrow = T, nrow = M) + error
predarray = array(NA, dim = c(M, length(x), nord))
for (i in 1:M) {
  for (j in 1:nord) {
```

```

    predarray[i, , j] = predict(lm(ymat[i, ] ~ poly(x, j, raw = TRUE)))
  }
}
# M matrices of size length(x) times nrd first, only look at
# variability in the M fits and plot M curves where we had 1 for
# plotting need to stack the matrices underneath eachother and make
# new variable 'rep'
stackmat = NULL
for (i in 1:M) {
  stackmat = rbind(stackmat, cbind(x, rep(i, length(x)), predarray[i,
    , ]))
}
# dim(stackmat)
colnames(stackmat) = c("x", "rep", paste("poly", 1:20, sep = ""))
sdf = as.data.frame(stackmat) #NB have poly1-20 now - but first only use 1,2,20
# to add true curve using stat_function - easiest solution
true_x = x
yrange = range(apply(sdf, 2, range)[, 3:22])
p1 = ggplot(data = sdf, aes(x = x, y = poly1, group = rep, colour = rep)) +
  scale_y_continuous(limits = yrange) + geom_line()
p1 = p1 + stat_function(fun = truefunc, lwd = 1.3, colour = "black") +
  ggtitle("poly1")
p2 = ggplot(data = sdf, aes(x = x, y = poly2, group = rep, colour = rep)) +
  scale_y_continuous(limits = yrange) + geom_line()
p2 = p2 + stat_function(fun = truefunc, lwd = 1.3, colour = "black") +
  ggtitle("poly2")
p10 = ggplot(data = sdf, aes(x = x, y = poly10, group = rep, colour = rep)) +
  scale_y_continuous(limits = yrange) + geom_line()
p10 = p10 + stat_function(fun = truefunc, lwd = 1.3, colour = "black") +
  ggtitle("poly10")
p20 = ggplot(data = sdf, aes(x = x, y = poly20, group = rep, colour = rep)) +
  scale_y_continuous(limits = yrange) + geom_line()
p20 = p20 + stat_function(fun = truefunc, lwd = 1.3, colour = "black") +
  ggtitle("poly20")
ggarrange(p1, p2, p10, p20)

```

```

## Warning: Multiple drawing groups in `geom_function()`. Did you use the correct
## `group`, `colour`, or `fill` aesthetics?

```

```

## Warning: Multiple drawing groups in `geom_function()`. Did you use the correct
## `group`, `colour`, or `fill` aesthetics?

```

```


## Warning: Multiple drawing groups in `geom_function()`. Did you use the correct
## `group`, `colour`, or `fill` aesthetics?

```

```


## Warning: Multiple drawing groups in `geom_function()`. Did you use the correct
## `group`, `colour`, or `fill` aesthetics?

```





Unsuprisingly, the second degree polynomial fits best.

b) The code is found in on the exercise sheet, We get the following plots.





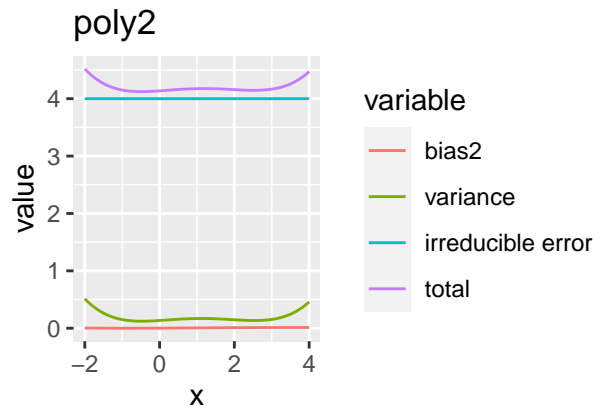

```
##
## Attaching package: 'reshape2'
## The following object is masked from 'package:tidyr':
##
##   smiths
```

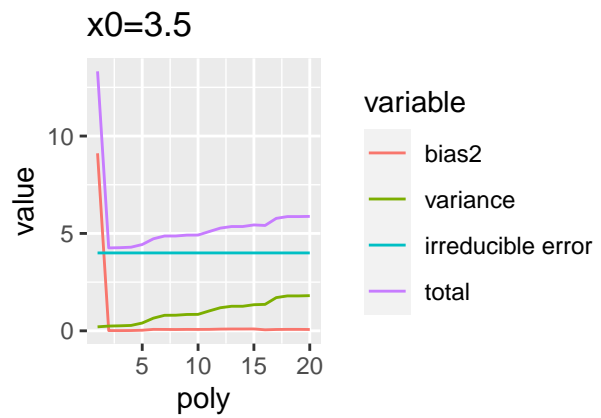





We can see that the second degree polynomial achieves the lowest mean testMSE, while the average trainMSE

increases with polynomial degree.

c)





d)